🎓 VIVA PRESENTATION GUIDE

Quantifying Customer Dissatisfaction — Hybrid NLP System

DSA Research Project | Final Viva Examination

> ***How to use this guide:*** *Read each section before you click through to that page. The "Say this to the panel" blocks are your speaking notes. The "Panel will ask" blocks are probable questions with model answers. Be confident — you built this!* 🚀

TABLE OF CONTENTS

#	Page / Phase	Key Concept to Explain
1	Home & Overview	Research question, 3-engine architecture
2	Data Hub	Multi-source input, Live Analysis mode
3	Phase I — Preprocessing	Text cleaning, feature engineering
4	Phase II — VADER Sentiment	Dissatisfaction Score formula
5	Phase II — LDA Topic Modeling	Latent theme discovery
6	Phase II — RoBERTa Sarcasm	Hybrid transformer + lexical boost
7	Phase III — Analytics Dashboard	Integrated insights, business value

---

🏆 YOUR RESEARCH IN ONE SENTENCE

> *"This system goes beyond simple positive/negative classification by quantifying **how dissatisfied** a customer is, **what topic** caused it, and **whether they were being sarcastic** — using a 3-phase hybrid NLP pipeline."*

PAGE 1 — 🏠 Home & Overview

What this page shows

- Project title, headline metrics (23,486 reviews, 6 topics, 3 engines)
- Three tabs: **Methodology**, **Research Phases**, **Dataset Info**

🗣️ Say this to the panel

> *"This is my project's landing page. The core research problem I'm solving is:* ***how do we move beyond a binary positive/negative label*** *and actually quantify customer dissatisfaction in a way that is actionable for a business?*
> *My answer is a* ***3-phase hybrid pipeline****: Phase I cleans and engineers the data, Phase II applies three complementary NLP engines, and Phase III presents integrated business intelligence. Each engine solves a different limitation of the others."*

The 3-Engine Table (explain this directly to the panel)

Engine	What it measures	Why it's needed
VADER	*Intensity* — how negative is the language?	Fast, rule-based, works on every review
LDA	*Theme* — what topic caused the complaint?	Tells the business *where* to fix things
RoBERTa	*Nuance* — is the customer being sarcastic?	Catches irony that fools VADER completely

❓ Panel Questions

Q: Why did you choose this dataset?

> *"The Women's Clothing E-Commerce dataset has 23,486 real customer reviews with star ratings, department labels, and age groups — making it ideal for multi-dimensional dissatisfaction analysis. The rating column also lets me validate my dissatisfaction scores against ground truth without needing manual labelling."*

Q: What is your research gap?

> *"Existing sentiment tools classify reviews as positive or negative. That tells a business nothing actionable. My system quantifies degree (0–100 scale), identifies root causes (6 topics), and detects hidden negativity (sarcasm) — three capabilities no single existing tool provides together."*

PAGE 2 — 📂 Data Hub

What this page shows

Three input modes: Built-in Dataset, Upload CSV/Excel, Live Review Text Analysis.

🗣️ Say this to the panel

> *"The Data Hub gives my system **three input modes** to make it production-ready, not just a research prototype. I'll demonstrate using the built-in Women's Clothing dataset. Click Load Women's Clothing Dataset — this loads all 23,486 reviews into memory."*

Demonstrate Live Analysis (very impressive for the panel)

1. Switch to ✏️ **Enter Live Review Text** mode
2. Load RoBERTa first (if not already loaded)
3. Paste this sarcastic review and click Analyse:

   > *"Oh, what a 'beautiful' surprise! I love how the buttons started falling off the > very first time I wore it. Truly a masterpiece of poor craftsmanship."*
4. Point to: VADER says **Positive** (compound +0.96) → RoBERTa detects **Sarcasm** → System overrides to **Negative (Sarcasm 🎭)** → Dissatisfaction Score: **95.6/100**

🗣️ Explain the live result

> *"This is the most powerful demonstration. VADER reads words like 'love' and 'beautiful' and gives a **+0.96 compound score** — nearly perfectly positive. But RoBERTa's bidirectional attention reads the full context and detects irony. My lexical boost layer adds +0.60 because it detects the quoted word 'beautiful', the phrase 'poor craftsmanship', and the rhetorical structure 'if you enjoy wasting money'. Combined irony score: **95.6%** → **Sarcasm confirmed** → **Sentiment corrected to Negative**."*

❓ Panel Questions

Q: Why support uploaded CSV files?

> *"To make the system generalizable. Any e-commerce company can upload their own review CSV (as long as it has a 'Review Text' column) and immediately get dissatisfaction analysis without changing a line of code."*

Q: How does the Live Analysis topic prediction work without LDA?

> *"For single reviews, training a full LDA model would be impractical. I implemented a keyword-based topic classifier using domain-specific keyword lists for all 6 topics. It counts keyword matches and assigns the topic with the highest score. This runs instantly and is accurate for clear-cut reviews."*

PAGE 3 — 🔧 Phase I: Preprocessing

What this page shows

- KPI row: rows after cleaning, rows removed, avg word count, negative reviews
- Tab 1 — Cleaned Data: side-by-side raw vs cleaned vs lemmatized text
- Tab 2 — Feature Stats: box plots of review*length, word*count, exclamation_count, etc.
- Tab 3 — Slang Examples: table of e-commerce slang normalisations

🗣️ Say this to the panel

> *"Phase I is the foundation of the entire pipeline. Raw user-generated text is noisy — it contains HTML tags, URLs, abbreviations like 'tts' (true to size), contractions like 'won't', and mixed-case inconsistencies. If we send dirty text to VADER or LDA, the results are unreliable.*
> *My preprocessing pipeline has **9 sequential steps**. Let me walk through them."*

The 9-Step Pipeline (point to Tab 3 while explaining)

Step	Operation	Why it matters
1	Drop duplicates & null reviews	Prevent model bias from repeated text
2	Strip HTML tags & URLs	Reviews scraped from web contain ` ` tags
3	Lowercase + remove non-ASCII	Standardise for lexicon lookup
4	Expand contractions	`won't` → `will not` (preserves negation for VADER)
5	Replace domain slang	`tts` → `true to size`, `gorge` → `gorgeous`
6	Remove special characters	Keep only letters and spaces
7	Tokenise + POS-tag	Identify noun/verb/adjective for correct lemmatisation
8	Lemmatise	`running` → `run`, `wore` → `wear`
9	Remove stopwords (preserve negation)	Remove `the`, `is`, but keep `not`, `never`

Key Technical Decision — Negation Preservation

> *"This is an important design choice I want to highlight. Standard stopword removal deletes words like 'not' and 'never'. But for sentiment analysis, removing 'not' would turn **'not good'** into **'good'** — completely reversing the meaning. I explicitly **preserve negation words** in my stopword filter."*

```
# From src/preprocessing.py
  NEGATION_WORDS = {'no', 'not', 'never', 'neither', 'nobody', 'nothing',
                    'nowhere', 'nor', "n't", 'cannot', 'without'}
  STOP_WORDS = set(stopwords.words('english')) - NEGATION_WORDS
```

Feature Engineering (show Tab 2)

Point to the box plots and explain:

> *"I engineer 7 derived features from the raw text. The panel can see that negative reviews tend to be **longer** (higher word count) and contain more **exclamation marks** — people who are angry write more. These features could be used as additional signals in a supervised classifier in future work."*

❓ Panel Questions

Q: What is lemmatisation and why use it over stemming?

> *"Lemmatisation uses vocabulary and morphological analysis to return a word to its **dictionary root form** — so 'wore', 'wearing', 'worn' all become 'wear'. Stemming just chops suffixes, producing non-words like 'wor'. For LDA topic modelling, lemmatisation produces much cleaner, interpretable topic words."*

Q: Why use POS-tagging before lemmatisation?

> *"The word 'wear' can be a noun or a verb. WordNet lemmatises them differently. Without POS tagging, 'better' (adjective) becomes 'better', but with the ADJ tag it correctly becomes 'good'. This improves accuracy especially for adjective-rich review text."*

Q: How many rows were removed in preprocessing?

> *"Approximately [check your KPI metric on screen]. These were removed due to: null review text, duplicated reviews, and reviews shorter than 10 characters (which are too short to carry meaningful sentiment)."*

PAGE 4 — 📊 Phase II: VADER Sentiment Analysis

What this page shows

- KPI row: Dissatisfaction Index, % Dissatisfied, % Severely Dissatisfied, Avg VADER Compound
- Dissatisfaction Gauge (0–100) + Sentiment Pie chart
- Dissatisfaction Histogram + Compound vs Rating Scatter
- Department-level dissatisfaction bar chart
- Review Explorer with severity filter
- Most dissatisfied review display

🗣️ Say this to the panel

> *"Phase II begins the actual NLP analysis. VADER — Valence Aware Dictionary and Sentiment Reasoner — is a rule-based sentiment analyser originally designed for social media. I've extended it with a **custom e-commerce lexicon** of 40+ terms with calibrated scores. For example, 'defective' = -3.5, 'overpriced' = -3.2, 'flattering' = +2.8, 'well-made' = +3.0."*

The Dissatisfaction Score Formula (write this on whiteboard if possible)

```
Dissatisfaction Score = max(0, −compound) × 100

  Where compound ∈ [−1.0, +1.0] from VADER

  Examples:
    compound = −0.80 → Score = 80/100  (Severely Dissatisfied)
    compound = −0.45 → Score = 45/100  (Highly Dissatisfied)
    compound = +0.60 → Score =  0/100  (Satisfied)
```

> *"I invert the compound score because VADER measures* ***sentiment*** *(positive = high), but I want to measure* ***dissatisfaction*** *(negative sentiment = high score). The max(0, ...) ensures satisfied reviews score 0, not a negative number."*

Severity Classification Thresholds

Score Range	Class	Business Action
70–100	🔴 Severely Dissatisfied	Immediate escalation required
45–69	☐ Highly Dissatisfied	Priority customer recovery
20–44	☐ Moderately Dissatisfied	Product improvement needed
5–19	☐ Mildly Dissatisfied	Monitor, low priority
0–4	✅ Satisfied	Positive customer experience

Walk the Panel Through Each Chart

Gauge Chart:

> *"The gauge shows our overall Dissatisfaction Index — the average dissatisfaction score across all reviews. For this dataset it's approximately [read from screen]. A score above 30 would be a business red flag."*

Compound vs Rating Scatter:

> *"This scatter plot validates my model. You can see a clear negative correlation — as star rating increases, VADER compound increases (more positive). This confirms that my NLP engine aligns with human star ratings, establishing model validity."*

Department Bar Chart:

> *"This shows which product department drives the most dissatisfaction. The panel can see [point to chart] that [e.g. Bottoms/Jackets] has the highest average dissatisfaction score — this is an actionable business insight."*

❓ Panel Questions

Q: Why VADER instead of TextBlob or a fine-tuned BERT?

> *"VADER has three advantages for this use case: First, it requires no training data — it works out-of-the-box on new domains. Second, it handles intensity modifiers like 'very bad' and 'extremely poor' correctly via its amplifier rules. Third, it runs in microseconds per review, making real-time analysis feasible. TextBlob is less accurate on informal text. Fine-tuned BERT would require labelled training data which I don't have for dissatisfaction scores."*

Q: What does the compound score mean exactly?

> *"The compound score is a normalised, weighted sum of all word valence scores in the lexicon, adjusted for rules like capitalisation (BAD scores more negative than bad), punctuation (!!! amplifies intensity), and degree modifiers (very, extremely, slightly). It ranges from -1.0 (most negative) to +1.0 (most positive)."*

Q: How did you calibrate your custom lexicon?

> *"I researched standard VADER lexicon scores for similar words and assigned values proportionally. For example, 'terrible' is -3.5 in VADER's standard lexicon, so I set domain-specific equivalents like 'defective' and 'unwearable' at similar levels based on how extreme they are in an e-commerce context."*

PAGE 5 — 🗂️ Phase II: LDA Topic Modeling

What this page shows

- KPI row: Topics Discovered (6), Coherence Score (C_v), Reviews Assigned, Dominant Topic
- Tab 1 — Topic Keywords: 6 cards showing top-8 keywords per topic with probabilities
- Tab 2 — Heatmap: Topic × Star Rating dissatisfaction pivot table
- Tab 3 — Topic Distribution: bar chart + pie chart of topic share
- Tab 4 — Labelled Reviews: filterable data table

🗣️ Say this to the panel

> *"VADER tells us how negative a review is. But it doesn't tell us why the customer is dissatisfied. That's what LDA Topic Modeling solves. LDA — Latent Dirichlet Allocation — is an unsupervised probabilistic model that discovers hidden thematic structure in text. It treats each review as a mixture of topics and each topic as a probability distribution over words."*

The 6 Topics Discovered

Topic ID	Label	Example Keywords
0	📏 Fit & Size	fit, size, runs, small, tight, chest, measurements
1	🏷️ Product Quality	fabric, material, quality, thin, stitching, faded, pilling
2	🚚 Delivery & Shipping	delivery, shipping, arrived, package, late, damaged
3	🤝 Customer Service	service, return, refund, exchange, support, response
4	💰 Value for Money	price, expensive, worth, overpriced, money, cost
5	🎨 Style & Design	style, design, beautiful, elegant, pattern, colour

Explain the Heatmap (Tab 2 — most powerful slide)

> *"This heatmap is perhaps the most actionable visualisation in the system. The rows are the 6 dissatisfaction topics. The columns are star ratings 1 to 5. Each cell is the **average dissatisfaction score** for reviews with that topic at that rating.*
> *Darker red = higher dissatisfaction. So you can see that [e.g. Product Quality] at 1-star has the highest average score — meaning quality complaints in 1-star reviews are the most severe. A business product manager should prioritise fixing exactly that intersection."*

LDA Hyperparameters — Justify Them

Parameter	Value	Justification
`num_topics`	6	Covers the main e-commerce complaint dimensions
`passes`	10	Sufficient corpus iterations for stable convergence
`alpha`	`auto`	Asymmetric Dirichlet — allows topics to have unequal prior probabilities
`eta`	`auto`	Word distribution prior — learned from data
`no_below`	5	Remove words appearing in fewer than 5 documents (noise)
`no_above`	0.6	Remove words in more than 60% of docs (too common, meaningless)

The Coherence Score (C_v)

> *"The coherence score measures how semantically similar the top words in each topic are. C_v ranges from 0 to 1. A score above 0.50 indicates meaningful topics. My model achieves [read from screen] — this confirms the 6 topics are distinct and interpretable, not random word clusters."*

❓ Panel Questions

Q: How does LDA work mathematically?

> *"LDA assumes a generative process: for each document, a topic distribution θ is sampled from a Dirichlet prior. For each word in the document, a topic z is sampled from θ, then a word w is drawn from that topic's word distribution φ_z. Training uses variational Bayes EM to find θ and φ that best explain the observed words. The output is: per-document topic proportions, and per-topic word distributions."*

Q: Why 6 topics and not 5 or 8?

> *"I chose 6 based on domain knowledge of e-commerce complaints (fit, quality, delivery, service, value, style are the standard complaint categories) and confirmed it with coherence score analysis. Fewer topics merged distinct themes; more topics created redundant, overlapping categories."*

Q: What is the difference between LDA and k-means clustering?

> *"K-means assigns each document to exactly one cluster. LDA assigns each document a probability distribution over all topics — a review can be 60% about Fit and 40% about Quality. This is more realistic because real reviews often discuss multiple themes. I use the dominant topic (highest probability) for assignment."*

PAGE 6 — 🎭 Phase II: RoBERTa Sarcasm Detection

What this page shows

- Model description and threshold (≥ 0.55)
- Sample size slider (100–3,000 reviews)
- KPI row: Reviews Analysed, Sarcastic Count, Sarcasm Rate %, High-Confidence Sarcasm
- Sarcasm Donut chart + Irony Probability Histogram
- Table of detected sarcastic reviews (sorted by irony_prob)
- VADER vs RoBERTa Disagreement Table

🗣 Say this to the panel

> *"This is the most technically advanced component of my system. The problem I'm solving is the **VADER-Sarcasm blind spot**: when a customer writes 'Oh great, the zipper broke after ONE wash — absolutely love it', VADER scores this **+0.78 positive** because it sees 'great' and 'love'. A human immediately recognises this as **deeply sarcastic**. My RoBERTa component catches exactly these cases."*

The RoBERTa Model

> *"I use `cardiffnlp/twitter-roberta-base-irony` — a RoBERTa transformer fine-tuned on Twitter irony detection. RoBERTa (Robustly Optimized BERT Pretraining Approach) uses bidirectional self-attention, meaning it reads the full sentence context in both directions simultaneously before making a prediction."*

The Hybrid Architecture — Most Important Technical Innovation

```
                Raw Review Text
                      │
          ┌───────────┴───────────┐
          │                       │
     RoBERTa Model          Lexical Boost
     (Context-based)        (Pattern-based)
     irony_prob = 0.356     boost = +0.600
          │                       │
          └───────────┬───────────┘
                      │
        combined = min(1.0, 0.356 + 0.600) = 0.956
                      │
      ≥ 0.55 threshold → IS SARCASTIC ✅
                      │
      Dissatisfaction Score = max(60, 95.6) = 95.6/100
```

Why a Lexical Boost Layer?

> *"The RoBERTa model was trained on **Twitter posts** — short, informal, emoji-heavy. E-commerce reviews are longer, more formal, and use different sarcasm patterns. This creates a **domain gap** — the model under-scores e-commerce sarcasm.*
> *My lexical boost layer compensates by detecting specific linguistic markers:"*

Pattern	Example	Boost
Quoted praise words	`'beautiful'` in quotes	+0.22
Explicit negative phrases	`poor craftsmanship`, `fell apart`	+0.12 each
Rhetorical structure	`if you enjoy wasting money`	+0.25
Classic opener	`Oh, what a beautiful surprise` + negative	+0.18
Contradictory superlative	`truly a masterpiece of poor...`	+0.22

> *"Crucially, the boost only activates when VADER's compound score is **above +0.3** — meaning VADER already thinks the review is positive. This prevents false positives on genuinely negative reviews."*

The VADER vs RoBERTa Disagreement Table

> *"This table at the bottom of the page is a research gold mine. It shows reviews where VADER compound > 0.1 (VADER thinks: positive/neutral) BUT RoBERTa detects sarcasm. These are the cases where a system using only VADER would have **completely mis-classified** the customer's true sentiment. My hybrid system catches them all."*

❓ Panel Questions

Q: Why not just use RoBERTa alone without VADER?

> *"Three reasons. First, RoBERTa is expensive — it takes seconds per review on CPU, making real-time analysis impractical for large datasets. VADER runs in microseconds. Second, RoBERTa is a classifier, not a scorer — it doesn't give a dissatisfaction intensity score (0–100). Third, the Twitter-trained model has domain gaps for e-commerce text that my hybrid lexical layer compensates for."*

Q: What does 'bidirectional attention' mean in RoBERTa?

> *"Traditional models like LSTMs read text left-to-right. BERT and RoBERTa use Transformer self-attention, allowing every word to attend to every other word simultaneously in both directions. When processing 'I absolutely love how it fell apart', RoBERTa can connect 'love' to 'fell apart' across the sentence, understanding the contradiction that indicates sarcasm."*

Q: How did you choose the 0.55 threshold?

> *"I tested multiple thresholds on my validation cases. At 0.50, genuine positive reviews occasionally get false-positive sarcasm flags. At 0.60, some clear sarcasm cases (like the 'masterpiece of poor craftsmanship' example) fall below the threshold. 0.55 provides the best precision-recall balance for this domain."*

Q: What is the sarcasm rate in your dataset?

> *"Approximately [read from your screen — typically 2–5%] of reviews show sarcasm. That may sound small, but in a dataset of 23,486 reviews, that's potentially 500–1,000 reviews being **wrongly classified as positive** by simple sentiment tools. For a business, those are customers they think are happy but are actually very dissatisfied."*

PAGE 7 — 📈 Phase III: Analytics Dashboard

What this page shows

- Executive Summary Banner (5 KPIs)
- Core Sentiment Overview: Gauge + Rating Distribution + Sentiment Pie
- Dissatisfaction Pain-Point Heatmap (Topic × Rating)
- Department & Demographic Analysis: Dept bar + Age box plot
- Word Clouds (Negative keywords / Positive keywords)
- Detailed Analysis: Compound vs Rating scatter + Topic bar
- Sarcasm Integration panel (if sarcasm has been run)
- Export section (Download CSV, Sarcasm CSV, KPI Summary)

🗣️ Say this to the panel

> *"Phase III is the business-facing output — the integrated analytics dashboard that synthesises all three NLP engines into one coherent view. I'll walk through each section and explain what business decisions each chart enables."*

Section-by-Section Walk-Through

Executive Summary Banner:

> *"Five headline KPIs at the top. The most important is the **Dissatisfaction Index** — a single number that tells a product manager how their customer base feels overall. This is what makes my system business-usable: decision-makers don't need to read 23,000 reviews, they look at one number."*

Dissatisfaction Gauge:

> *"The gauge visually represents the index on a 0–100 arc. Anything above 30 should be a concern for the business. The dial position immediately communicates severity without the viewer needing to interpret numbers."*

Rating Distribution Bar Chart:

> *"This shows the star rating breakdown. For this dataset, 55.9% of reviews are 5-star. This confirms what academic literature calls **review positivity bias** — customers who are very happy or very unhappy are most motivated to write reviews."*

Topic × Rating Heatmap:

> *"As I explained in the LDA section, this is the most actionable chart. Cross-referencing topics with ratings tells a business which specific complaint category is most severe. For example, if Delivery issues at 1-star are brightest red, the business should prioritise logistics improvements."*

Word Clouds:

> *"Two word clouds — negative reviews (1–2 stars) and positive reviews (4–5 stars). The negative cloud reveals the actual vocabulary customers use when unhappy. This is useful for marketing teams to understand brand perception and for product teams to identify recurring failure keywords."*

Age Box Plot:

> *"This shows dissatisfaction score distribution across age groups. Older customers may have different expectations. If a particular age group shows higher dissatisfaction, targeted communication strategies can be developed."*

Sarcasm Integration Panel:

> *"This panel only appears if you've run the RoBERTa sarcasm detection. It overlays the sarcasm findings onto the dashboard and highlights the key research insight: sarcastic reviews carry hidden dissatisfaction that inflates the apparent satisfaction rate if not detected."*

Export Section:

> *"The system produces three downloadable outputs: the full analysed dataset with all NLP columns appended, the sarcasm-specific results, and a KPI summary CSV. This makes the system integration-ready for business BI tools like Power BI or Tableau."*

❓ Panel Questions

Q: How does the dashboard handle missing data? (e.g., if sarcasm hasn't been run)

> *"Every section is conditional. The sarcasm panel only renders if `sarcasm_done` is True in session state. Charts check whether required columns exist before rendering. If a department or age column is missing from an uploaded CSV, those charts are gracefully skipped with an informative message."*

Q: What business decisions can a company make from this dashboard?

> *"Five concrete decisions: 1) If Dissatisfaction Index > 30, launch a customer recovery programme. 2) The heatmap identifies which product category and price point to improve first. 3) The department bar shows which team to hold accountable. 4) The sarcasm rate shows how much of their '5-star' base is actually dissatisfied. 5) The word clouds guide copywriting and response templates."*

🔧 TECHNICAL ARCHITECTURE SUMMARY

File Structure (say this if asked)

```
project/
  ├── app.py                    # Streamlit UI — 7 pages, all logic
  ├── src/
  │   ├── preprocessing.py      # Phase I: text cleaning, lemmatisation, feature engine
  ering
  │   ├── sentiment_vader.py    # Phase II-A: VADER + custom lexicon + dissatisfaction
  score
  │   ├── topic_modeling.py     # Phase II-B: LDA training, topic assignment, coherence
  │   ├── sarcasm_detector.py   # Phase II-C: RoBERTa + lexical boost hybrid
  │   └── visualizations.py     # All Plotly charts (gauge, heatmap, donut, scatter…)
  └── dataset/
      └── Womens Clothing E-Commerce Reviews.csv
```

Data Flow Through the Pipeline

```
Raw CSV (23,486 reviews)
        ↓
  [Phase I] preprocessing.py
    → cleaned_text, processed_text, word_count, exclamation_count, sentiment_label
        ↓
  [Phase II-A] sentiment_vader.py
    → compound, vader_pos, vader_neg, dissatisfaction_score, sentiment_class
        ↓
  [Phase II-B] topic_modeling.py
    → dominant_topic_id, topic_label, topic_probability
        ↓
  [Phase II-C] sarcasm_detector.py
    → irony_prob, roberta_irony, lexical_boost, is_sarcastic
        ↓
  [Phase III] Dashboard + Export
    → Integrated business intelligence visualisations
```

Session State Architecture

> *"I use Streamlit's session_state to persist processed DataFrames between page navigations. This means heavy computations (VADER on 23K reviews, LDA training, RoBERTa inference) only run once per session, not on every page reload. The pipeline status flags (preprocessingdone, vaderdone, topicdone, sarcasmdone) enforce sequential execution order."*

❓ RAPID-FIRE PANEL QUESTIONS & ANSWERS

Q: What is the overall accuracy of your system?

> *"My validation approach is indirect but meaningful. I compare VADER compound scores against ground-truth star ratings — the negative correlation ($r \approx -0.65$) confirms the system's directional accuracy. For sarcasm, I validated on hand-crafted test cases covering all major irony patterns. A formal precision/recall study on a labelled sarcasm subset would be the logical next step."*

Q: Can this system handle languages other than English?

> *"Currently no — VADER's lexicon is English-only, and the RoBERTa model was trained on English Twitter data. Extension to multilingual support would require a multilingual VADER equivalent and a multilingual irony detection model, such as XLM-RoBERTa."*

Q: What are the limitations of your system?

> *"Three main limitations: 1) VADER struggles with highly contextual irony without the lexical boost layer. 2) LDA is an unsupervised method — the 6 topics were pre-labelled by me based on domain knowledge; a different analyst might label them differently. 3) The RoBERTa model has a domain gap (Twitter → E-commerce) that my lexical boost partially compensates for, but edge cases remain."*

Q: How would you improve this system?

> *"Four directions: 1) Fine-tune RoBERTa on e-commerce sarcasm data for better domain fit. 2) Add a supervised classification layer using the engineered features (word count, exclamation ratio) as additional signals. 3) Replace keyword-based live topic classification with a pre-trained zero-shot classifier. 4) Add temporal analysis — tracking how the dissatisfaction index changes month-over-month."*

Q: What makes this a research contribution vs. just a dashboard?

> *"Three contributions: 1) The novel **hybrid sarcasm detection architecture** combining a transformer model with a domain-specific lexical boost layer — this addresses a gap in existing e-commerce sentiment literature. 2) The **Dissatisfaction Index** formulation (0–100 scale from VADER compound) as an actionable business metric. 3) The **integration of three complementary NLP techniques** in a single pipeline where each engine compensates for the others' weaknesses."*

🚀 PRESENTATION FLOW CHECKLIST

Before your viva, run through this in order:

- [ ] Launch app: `python -m streamlit run app.py`
- [ ] **Page 1** (Home): Explain the 3-engine table — 2 minutes
- [ ] **Page 2** (Data Hub): Load dataset + Demo live sarcasm analysis — 3 minutes
- [ ] **Page 3** (Preprocessing): Show cleaned text table + explain 9 steps — 2 minutes
- [ ] **Page 4** (VADER): Show gauge + scatter + write formula on whiteboard — 3 minutes
- [ ] **Page 5** (LDA): Show keyword cards + heatmap (most impressive chart) — 3 minutes
- [ ] **Page 6** (RoBERTa): Explain hybrid architecture + show disagreement table — 4 minutes
- [ ] **Page 7** (Dashboard): Show executive summary + word clouds — 2 minutes
- [ ] **Q&A**: Use answers from this guide — 10 minutes

Total demo time: ~20 minutes + Q&A

💡 FINAL CONFIDENCE TIPS

1. **Start with the live demo** — paste the sarcastic review in Data Hub first. It's the most dramatic demonstration and immediately shows the system works.
2. **Always connect to business value** — after every technical explanation, say *"and this means a business can..."* — panels love practical impact.
3. **The heatmap is your star slide** — spend extra time on it. It's visual, intuitive, and directly actionable.
4. **Own your design choices** — if they ask "why VADER?", don't say "because I found it online." Say "because it meets these specific requirements: speed, no training data needed, and rules-based interpretability."
5. **Numbers matter** — know these by heart:
   - **23,486** reviews analysed
   - **6** dissatisfaction topics
   - **3** NLP engines
   - **0.55** sarcasm threshold
   - **40+** custom lexicon entries
   - **9** preprocessing steps

Guide prepared for: Quantifying Customer Dissatisfaction — DSA Research Project System built with: Python · Streamlit · NLTK · Gensim · HuggingFace Transformers · Plotly